

EDUCATION

University of Toronto

MSc in Applied Computing (Computer Science)

Courses: Computational Imaging, Imitation Learning for Robotics, Natural Language Computing, Visual Computing Systems

Sep 2022 – Dec 2023 (expected)

GPA: 4.0 / 4.0

Indian Institute of Technology Madras

BTech + MTech in Aerospace Engineering and Data Science

Courses: Computational Photography, Deep Learning, Systems Engineering for Deep Learning

Jul 2022

GPA: 9.26 / 10.00

EXPERIENCE

Advanced Micro Devices (AMD)

Applied Research Intern, Machine Learning and 3D Computer Vision

May 2023 - Dec 2023

- Built an automatic system to produce high-fidelity animatable 3D human avatars from a consumer laptop webcam.
- Developed a novel method based on photorealistic GANs, 3D morphable models, landmark detection, face recognition embeddings and differentiable rendering in PyTorch3D.
- Achieved ~45% higher perceptual quality and ~25% better lighting symmetry compared to existing literature when applied to low-quality webcams, while also satisfying latency constraints for edge devices.
- Contributed to a suite of apps running on the AI inference engine in the next generation of Ryzen mobile chips.

Indian Institute of Technology Madras

Student Researcher, Deep Learning for Fluid Mechanics and CFD

Jul 2020 - Jul 2022

- Worked on developing fast and accurate AI algorithms for fluid simulations.
- Trained a ConvNet with physics-based components in JAX (framework) to correct errors in low-precision simulations.
- Published two first-author research papers and presented at a conference.
- Demonstrated improved computational efficiency and ~45% error reduction: [\[Paper\]](#)

unMazer.ai

Intern, Computer Vision and Product Development

Summer 2020

- Event detection using a custom CNN+LSTM network and object detection using YOLOv4 in CCTV road footage.
- Built a proctoring tool to detect cheating in online exams using eye tracking models and OpenCV.
- Instrumental in building three AI-based minimum viable products (MVPs) during the company's initial phase.
- Blog post: [\[Leveraging AI to Reduce Road Fatalities\]](#)

SELECTED PROJECTS

LLM Learning Assistant: Long-Form Content Summarization and Q&A using Retrieval-Augmented Generation [\[Code\]](#)

Built an app to help you quickly comprehend the content of podcasts and lectures, capable of running locally via llama.cpp.

Learned to work with Llama 2, LangChain, and vector databases to perform search and retrieval-augmented generation.

Deconvolution using ADMM with Diffusion Denoising Prior [\[Poster\]](#) [\[Paper\]](#) [\[Code\]](#)

Improved image deblurring under high Gaussian noise by using a pretrained diffusion model as a denoising prior.

Fall 2022

DADAgger: Imitation Learning with Disagreement-Augmented Dataset Aggregation [\[Paper\]](#) [\[Code\]](#)

Improved data efficiency of the DAgger algorithm by estimating model uncertainty using ensemble variance.

Resulted in a car-racing agent that learns successfully using 80% fewer queries to the expert.

Fall 2022

Find more course projects and side-projects at <https://akasharidas.github.io/projects/>

SELECTED PUBLICATIONS

Akash Haridas, Nagabhushana Rao Vadlamani, Yuki Minamoto, [Deep neural networks to correct sub-precision errors in CFD](#),

Applications in Energy and Combustion Science, Volume 12, 2022

SKILLS

Programming: Python, C++, Julia, MATLAB, CUDA, HTML/CSS, LaTeX

Tools/Frameworks: PyTorch, LangChain, JAX, NumPy, Linux, OpenCV, Pandas, Scikit Learn, Google Cloud Platform, PySpark